

Notebook

March 10, 2023

1 Student Grade Analysis & Prediction

Objective: Prediction of the final grade of Portugese high school students

Data Set Information The data used is from a Portuguese secondary school. The data includes academic and personal characteristics of the students as well as final grades. The task is to predict the final grade from the student information. (Regression) [Link to dataset](#)

Citation: P. Cortez and A. Silva. Using Data Mining to Predict Secondary School Student Performance. In A. Brito and J. Teixeira Eds., Proceedings of 5th FUTURE BUSINESS TECHNOLOGY Conference (FUBUTEC 2008) pp. 5-12, Porto, Portugal, April, 2008, EUROSIS, ISBN 978-9077381-39-7. [Web Link](#)

1.1 Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

1.2 The Data

Let's start by reading in the student-mat.csv file into a pandas dataframe.

```
stud= pd.read_csv("gs://iitj_bigdatapproject_pgd_dev/Pub/student-mat.csv/*.csv")
→ # Read the dataset
```

```
print('Total number of students:',len(stud))
```

Total number of students: 395

```
stud['G3'].describe()
```

```
count    395.000000
mean      10.415190
std        4.581443
min        0.000000
```

```
25%      8.000000
50%     11.000000
75%     14.000000
max      20.000000
Name: G3, dtype: float64
```

```
stud.info()    # Information on dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
school      395 non-null object
sex         395 non-null object
age         395 non-null int64
address     395 non-null object
famsize     395 non-null object
Pstatus     395 non-null object
Medu        395 non-null int64
Fedu        395 non-null int64
Mjob        395 non-null object
Fjob        395 non-null object
reason      395 non-null object
guardian    395 non-null object
traveltime  395 non-null int64
studytime   395 non-null int64
failures    395 non-null int64
schoolsup   395 non-null object
famsup      395 non-null object
paid        395 non-null object
activities  395 non-null object
nursery     395 non-null object
higher      395 non-null object
internet    395 non-null object
romantic    395 non-null object
famrel      395 non-null int64
freetime    395 non-null int64
goout       395 non-null int64
Dalc        395 non-null int64
Walc        395 non-null int64
health      395 non-null int64
absences    395 non-null int64
G1          395 non-null int64
G2          395 non-null int64
G3          395 non-null int64
dtypes: int64(16), object(17)
memory usage: 101.9+ KB
```

```
stud.columns    # Dataset Columns
```

```
Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
      'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
      'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
      'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
      'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
      dtype='object')
```

```
stud.describe()    # Dataset description
```

	age	Medu	Fedu	traveltime	studytime	failures	\
count	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	
mean	16.696203	2.749367	2.521519	1.448101	2.035443	0.334177	
std	1.276043	1.094735	1.088201	0.697505	0.839240	0.743651	
min	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	
25%	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	
50%	17.000000	3.000000	2.000000	1.000000	2.000000	0.000000	
75%	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	
max	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	

	famrel	freetime	goout	Dalc	Walc	health	\
count	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	
mean	3.944304	3.235443	3.108861	1.481013	2.291139	3.554430	
std	0.896659	0.998862	1.113278	0.890741	1.287897	1.390303	
min	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	4.000000	3.000000	2.000000	1.000000	1.000000	3.000000	
50%	4.000000	3.000000	3.000000	1.000000	2.000000	4.000000	
75%	5.000000	4.000000	4.000000	2.000000	3.000000	5.000000	
max	5.000000	5.000000	5.000000	5.000000	5.000000	5.000000	

	absences	G1	G2	G3
count	395.000000	395.000000	395.000000	395.000000
mean	5.708861	10.908861	10.713924	10.415190
std	8.003096	3.319195	3.761505	4.581443
min	0.000000	3.000000	0.000000	0.000000
25%	0.000000	8.000000	9.000000	8.000000
50%	4.000000	11.000000	11.000000	11.000000
75%	8.000000	13.000000	13.000000	14.000000
max	75.000000	19.000000	19.000000	20.000000

```
stud.head()    # First 5 values of dataset
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	

2	GP	F	15	U	LE3	T	1	1	at_home	other	...
3	GP	F	15	U	GT3	T	4	2	health	services	...
4	GP	F	16	U	GT3	T	3	3	other	other	...

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	6	5	6	6
1	5	3	3	1	1	3	4	5	5	6
2	4	3	2	2	3	3	10	7	8	10
3	3	2	2	1	1	5	2	15	14	15
4	4	3	2	1	2	5	4	6	10	10

[5 rows x 33 columns]

```
stud.tail()      # Last 5 values of dataset
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
390	MS	M	20	U	LE3	A	2	2	services	services	
391	MS	M	17	U	LE3	T	3	1	services	services	
392	MS	M	21	R	GT3	T	1	1	other	other	
393	MS	M	18	R	LE3	T	3	2	services	other	
394	MS	M	19	U	LE3	T	1	1	other	at_home	

	...	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
390	...	5	5	4	4	5	4	11	9	9	9
391	...	2	4	5	3	4	2	3	14	16	16
392	...	5	5	3	3	3	3	3	10	8	7
393	...	4	4	1	3	4	5	0	11	12	10
394	...	3	2	3	3	3	5	5	8	9	9

[5 rows x 33 columns]

```
stud.isnull().any()      # To check any null values present in dataset
```

school	False
sex	False
age	False
address	False
famsize	False
Pstatus	False
Medu	False
Fedu	False
Mjob	False
Fjob	False
reason	False
guardian	False
traveltime	False

```

studytime      False
failures       False
schoolsup      False
famsup         False
paid           False
activities     False
nursery        False
higher         False
internet       False
romantic       False
famrel         False
freetime       False
goout          False
Dalc           False
Walc           False
health         False
absences       False
G1             False
G2             False
G3             False
dtype: bool

```

```

import cufflinks as cf
cf.go_offline()

```

```

stud.iplot()      # Plot for the all attributes

```

```

stud.iplot(kind='scatter',x='age',y='G3',mode='markers',size=8)      # Plot for age vs G3

```

```

stud.iplot(kind='box')

```

```

stud['G3'].iplot(kind='hist',bins=100,color='blue')

```

2 Data Visualization

```


sns.heatmap(stud.isnull(),cmap="rainbow",yticklabels=False)      # To check any null values present in dataset pictorially

```

```


<matplotlib.axes._subplots.AxesSubplot at 0x2041744c438>

```



```
sns.heatmap(stud.isnull(),cmap="viridis",yticklabels=False) # Map color ->
    ↪ viridis
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x20414b1c2b0>
```



- There are no null values in the given dataset

2.1 Student's Sex


```
f_stud = len(stud[stud['sex'] == 'F'])    # Number of female students
print('Number of female students:',f_stud)
m_stud = len(stud[stud['sex'] == 'M'])    # Number of male students
print('Number of male students:',m_stud)
```

Number of female students: 208

Number of male students: 187

```
sns.set_style('whitegrid')    # male & female student representaion on countplot
sns.countplot(x='sex',data=stud,palette='plasma')
```


<matplotlib.axes._subplots.AxesSubplot at 0x204149de630>




- The gender distribution is pretty even.

3 Age of Students

```
b = sns.kdeplot(stud['age'])    # Kernel Density Estimations
b.axes.set_title('Ages of students')
b.set_xlabel('Age')
b.set_ylabel('Count')
plt.show()
```

```
b = sns.countplot(x='age',hue='sex', data=stud, palette='inferno')
b.axes.set_title('Number of Male & Female students in different age groups')
b.set_xlabel("Age")
b.set_ylabel("Count")
plt.show()
```



- The student age seems to be ranging from 15-19, where gender distribution is pretty even in each age group.
- The age group above 19 may be outliers, year back students or droupouts.


4 Students from Urban & Rural Areas

```
u_stud = len(stud[stud['address'] == 'U'])    # Number of urban areas students
print('Number of Urban students:',u_stud)
r_stud = len(stud[stud['address'] == 'R'])    # Number of rural areas students
print('Number of Rural students:',r_stud)
```

Number of Urban students: 307
Number of Rural students: 88

```
sns.set_style('whitegrid')
sns.countplot(x='address',data=stud,palette='magma')    # urban & rural
→representaion on countplot
```


<matplotlib.axes._subplots.AxesSubplot at 0x2041470feb8>



- Approximately 77.72% students come from urban region and 22.28% from rural region.

```
sns.countplot(x='address',hue='G3',data=stud,palette='Oranges')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x204158e1710>
```




5 EDA - Exploratory Data Analysis

5.0.1 1. Does age affect final grade?

```
b= sns.boxplot(x='age', y='G3',data=stud,palette='gist_heat')
b.axes.set_title('Age vs Final Grade')
```


```
Text(0.5, 1.0, 'Age vs Final Grade')
```



- Plotting the distribution rather than statistics would help us better understand the data.
- The above plot shows that the median grades of the three age groups(15,16,17) are similar. Note the skewness of age group 19. (may be due to sample size). Age group 20 seems to score highest grades among all.


```
b = sns.swarmplot(x='age', y='G3', hue='sex', data=stud, palette='PiYG')
b.axes.set_title('Does age affect final grade?')
```

```
Text(0.5, 1.0, 'Does age affect final grade?')
```



5.1 2. Do urban students perform better than rural students?

```
# Grade distribution by address
sns.kdeplot(stud.loc[stud['address'] == 'U', 'G3'], label='Urban', shade = True)
sns.kdeplot(stud.loc[stud['address'] == 'R', 'G3'], label='Rural', shade = True)
plt.title('Do urban students score higher than rural students?')
plt.xlabel('Grade')
plt.ylabel('Density')
plt.show()
```



- The above graph clearly shows there is not much difference between the grades based on location.

```
stud.corr()['G3'].sort_values()
```

```
failures    -0.360415
age         -0.161579
goout       -0.132791
traveltime  -0.117142
health      -0.061335
Dalc        -0.054660
Walc        -0.051939
freetime    0.011307
absences    0.034247
famrel      0.051363
studytime   0.097820
Fedu        0.152457
Medu        0.217147
G1          0.801468
G2          0.904868
G3          1.000000
Name: G3, dtype: float64
```

5.2 Encoding categorical variables using LabelEncoder()

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
stud.iloc[:,0]=le.fit_transform(stud.iloc[:,0])
stud.iloc[:,1]=le.fit_transform(stud.iloc[:,1])
stud.iloc[:,3]=le.fit_transform(stud.iloc[:,3])
stud.iloc[:,4]=le.fit_transform(stud.iloc[:,4])
stud.iloc[:,5]=le.fit_transform(stud.iloc[:,5])
stud.iloc[:,8]=le.fit_transform(stud.iloc[:,8])
stud.iloc[:,9]=le.fit_transform(stud.iloc[:,9])
stud.iloc[:,10]=le.fit_transform(stud.iloc[:,10])
stud.iloc[:,11]=le.fit_transform(stud.iloc[:,11])
stud.iloc[:,15]=le.fit_transform(stud.iloc[:,15])
stud.iloc[:,16]=le.fit_transform(stud.iloc[:,16])
stud.iloc[:,17]=le.fit_transform(stud.iloc[:,17])
stud.iloc[:,18]=le.fit_transform(stud.iloc[:,18])
stud.iloc[:,19]=le.fit_transform(stud.iloc[:,19])
stud.iloc[:,20]=le.fit_transform(stud.iloc[:,20])
stud.iloc[:,21]=le.fit_transform(stud.iloc[:,21])
stud.iloc[:,22]=le.fit_transform(stud.iloc[:,22])
```

```
stud.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	\
0	0	0	18	1	0	0	4	4	0	4
1	0	0	17	1	0	1	1	1	0	2
2	0	0	15	1	1	1	1	1	0	2
3	0	0	15	1	0	1	4	2	1	3
4	0	0	16	1	0	1	3	3	2	2

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	6	5	6	6
1	5	3	3	1	1	3	4	5	5	6
2	4	3	2	2	3	3	10	7	8	10
3	3	2	2	1	1	5	2	15	14	15
4	4	3	2	1	2	5	4	6	10	10

[5 rows x 33 columns]

```
stud.tail()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	\
390	1	1	20	1	1	0	2	2	3	3
391	1	1	17	1	1	1	3	1	3	3
392	1	1	21	0	0	1	1	1	2	2
393	1	1	18	0	1	1	3	2	3	2

394	1	1	19	1	1	1	1	1	2	0	...
	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3	
390	5	5	4	4	5	4	11	9	9	9	
391	2	4	5	3	4	2	3	14	16	16	
392	5	5	3	3	3	3	3	10	8	7	
393	4	4	1	3	4	5	0	11	12	10	
394	3	2	3	3	3	5	5	8	9	9	

[5 rows x 33 columns]

```
stud.corr()['G3'].sort_values()      # Correlation wrt G3
```

```
failures      -0.360415
age           -0.161579
goout         -0.132791
romantic      -0.129970
traveltime    -0.117142
schoolsup     -0.082788
guardian      -0.070109
health        -0.061335
Pstatus       -0.058009
Dalc          -0.054660
Walc          -0.051939
school        -0.045017
famsup        -0.039157
freetime      0.011307
activities    0.016100
absences      0.034247
Fjob          0.042286
famrel        0.051363
nursery       0.051568
famsize       0.081407
studytime     0.097820
internet      0.098483
paid          0.101996
Mjob          0.102082
sex           0.103456
address       0.105756
reason        0.121994
Fedu          0.152457
higher        0.182465
Medu          0.217147
G1            0.801468
G2            0.904868
G3            1.000000
Name: G3, dtype: float64
```

```
# drop the school and grade columns
stud = stud.drop(['school', 'G1', 'G2'], axis='columns')
```

- Although G1 and G2 which are period grades of a student and are highly correlated to the final grade G3, we drop them. It is more difficult to predict G3 without G2 and G1, but such prediction is much more useful because we want to find other factors affect the grade.

```
# Find correlations with the Grade
most_correlated = stud.corr().abs()['G3'].sort_values(ascending=False)

# Maintain the top 8 most correlation features with Grade
most_correlated = most_correlated[:9]
most_correlated
```

```
G3          1.000000
failures    0.360415
Medu        0.217147
higher      0.182465
age         0.161579
Fedu        0.152457
goout       0.132791
romantic    0.129970
reason      0.121994
Name: G3, dtype: float64
```


```
stud = stud.loc[:, most_correlated.index]
stud.head()
```

	G3	failures	Medu	higher	age	Fedu	goout	romantic	reason
0	6	0	4	1	18	4	4	0	0
1	6	0	1	1	17	1	3	0	0
2	10	3	1	1	15	1	2	0	2
3	15	0	4	1	15	2	2	1	1
4	10	0	3	1	16	3	2	0	1

5.2.1 Failure Attribute

```
b = sns.swarmplot(x=stud['failures'],y=stud['G3'],palette='autumn')
b.axes.set_title('Previous Failures vs Final Grade(G3)')
```

```
Text(0.5, 1.0, 'Previous Failures vs Final Grade(G3)')
```




Observation : Student with less previous failures usually score higher

5.2.2 Family Education Attribute (Fedu + Medu)

```
fa_edu = stud['Fedu'] + stud['Medu']
b = sns.swarmplot(x=fa_edu,y=stud['G3'],palette='summer')
b.axes.set_title('Family Education vs Final Grade(G3)')
```

```
Text(0.5, 1.0, 'Family Education vs Final Grade(G3)')
```




Observation : Educated families result in higher grades

5.2.3 Wish to go for Higher Education Attribute

```
b = sns.boxplot(x=stud['higher'],y=stud['G3'],palette='binary')
b.axes.set_title('Higher Education vs Final Grade(G3)')
```

```
Text(0.5, 1.0, 'Higher Education vs Final Grade(G3)')
```

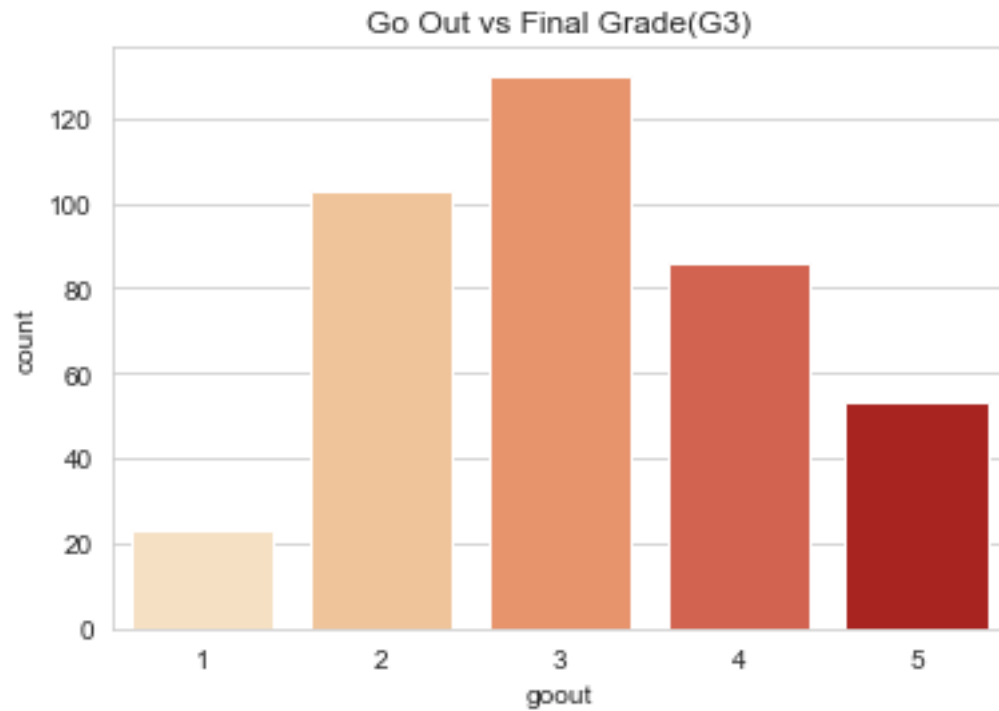


Observation : Students who wish to go for higher studies score more

5.3 Going Out with Friends Attribute

```
b = sns.countplot(x=stud['goout'],palette='OrRd')
b.axes.set_title('Go Out vs Final Grade(G3)')
```

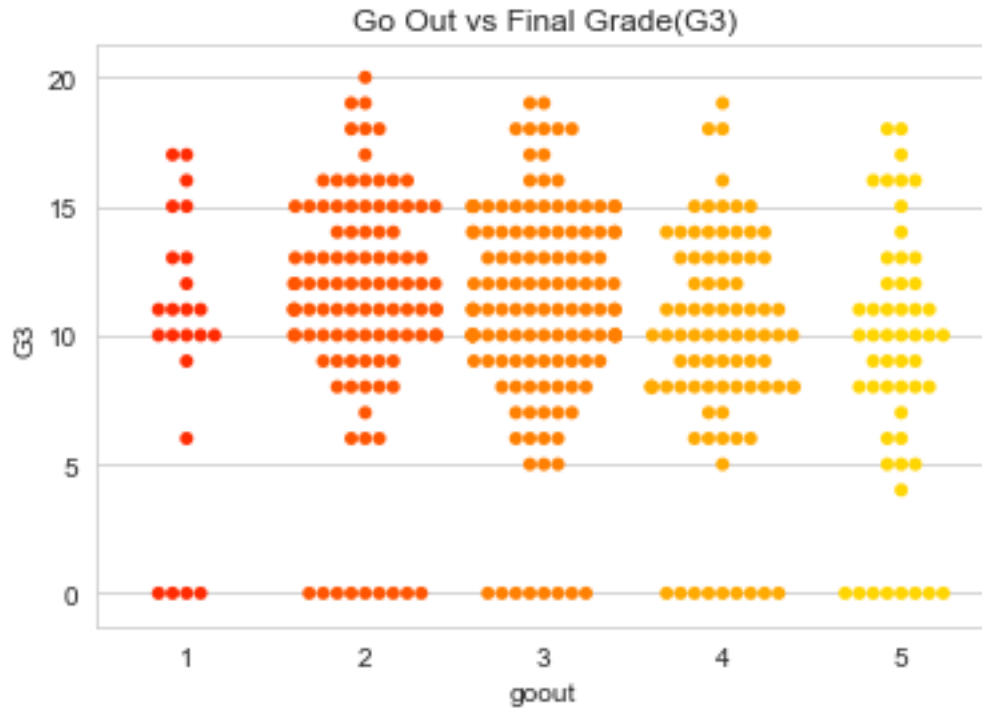
```
Text(0.5, 1.0, 'Go Out vs Final Grade(G3)')
```



Observation : The students have an average score when it comes to going out with friends.

```
b = sns.swarmplot(x=stud['goout'],y=stud['G3'],palette='autumn')  
b.axes.set_title('Go Out vs Final Grade(G3)')
```

```
Text(0.5, 1.0, 'Go Out vs Final Grade(G3)')
```

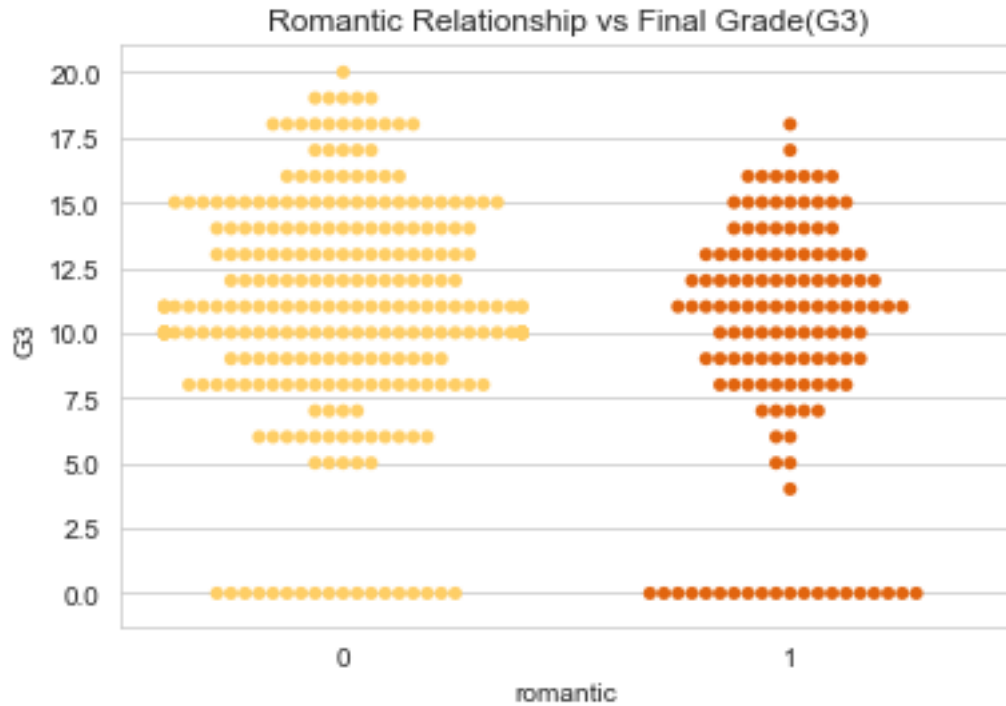


Observation : Students who go out a lot score less

5.3.1 Romantic relationship Attribute

```
b = sns.swarmplot(x=stud['romantic'],y=stud['G3'],palette='YlOrBr')
b.axes.set_title('Romantic Relationship vs Final Grade(G3)')
```

```
Text(0.5, 1.0, 'Romantic Relationship vs Final Grade(G3)')
```



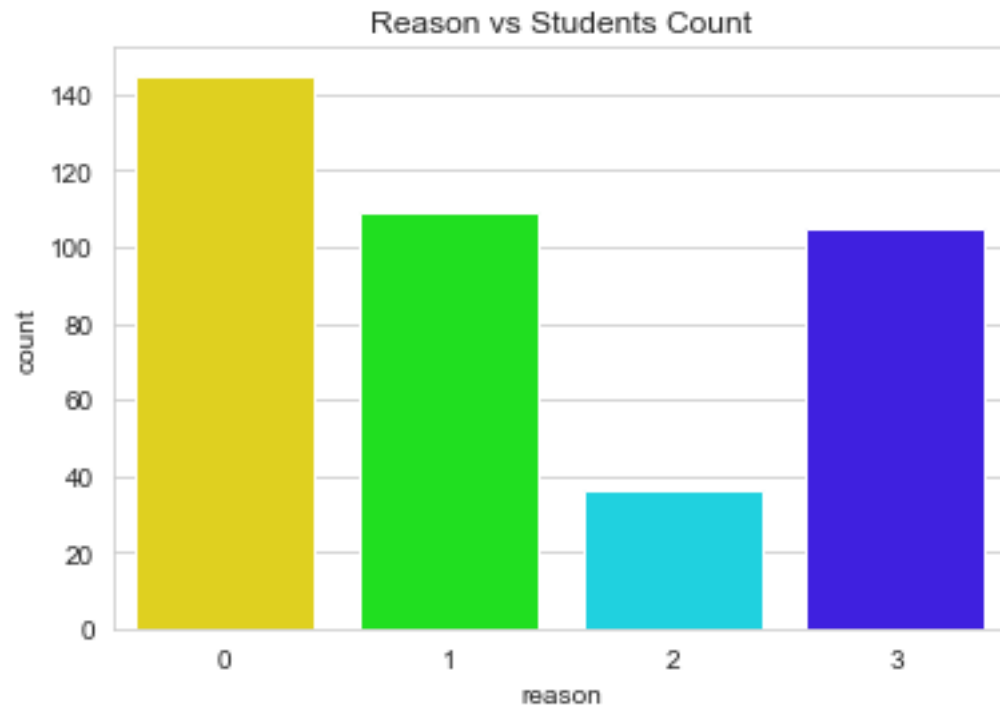
- Here romantic attribute with value 0 means no relationship and value with 1 means in relationship

Observation : Students with no romantic relationship score higher

5.3.2 Reason Attribute


```
b = sns.countplot(x='reason',data=stud,palette='gist_rainbow')    # Reason to
    ↪choose this school
b.axes.set_title('Reason vs Students Count')
```

```
Text(0.5, 1.0, 'Reason vs Students Count')
```

```
b = sns.swarmplot(x='reason', y='G3', data=stud,palette='gist_rainbow')  
b.axes.set_title('Reason vs Final grade')
```

```
Text(0.5, 1.0, 'Reason vs Final grade')
```



Observation : The students have an equally distributed average score when it comes to reason attribute.

6 Machine Learning Algorithms

```
# Standard ML Models for comparison
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVR

# Splitting data into training/testing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

# Metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, median_absolute_error

# Distributions
import scipy
```

```
# splitting the data into training and testing data (75% and 25%)
# we mention the random state to achieve the same split everytime we run the
→code
X_train, X_test, y_train, y_test = train_test_split(stud, stud['G3'], test_size=
→= 0.25, random_state=42)
```

```
X_train.head()
```

	G3	failures	Medu	higher	age	Fedu	goout	romantic	reason
16	14	0	4	1	16	4	3	0	3
66	12	0	4	1	15	4	3	1	3
211	13	0	4	1	17	4	5	1	1
7	6	0	4	1	17	4	4	0	1
19	10	0	4	1	16	3	3	0	1

6.1 MAE - Mean Absolute Error & RMSE - Root Mean Square Error

```
# Calculate mae and rmse
def evaluate_predictions(predictions, true):
    mae = np.mean(abs(predictions - true))
    rmse = np.sqrt(np.mean((predictions - true) ** 2))

    return mae, rmse
```

```
# find the median
median_pred = X_train['G3'].median()

# create a list with all values as median
median_preds = [median_pred for _ in range(len(X_test))]

# store the true G3 values for passing into the function
true = X_test['G3']
```

```
# Display the naive baseline metrics
mb_mae, mb_rmse = evaluate_predictions(median_preds, true)
print('Median Baseline MAE: {:.4f}'.format(mb_mae))
print('Median Baseline RMSE: {:.4f}'.format(mb_rmse))
```

Median Baseline MAE: 3.7879

Median Baseline RMSE: 4.8252

```
# Evaluate several ml models by training on training set and testing on testing
→set
def evaluate(X_train, X_test, y_train, y_test):
    # Names of models
    model_name_list = ['Linear Regression', 'ElasticNet Regression',
```

```

        'Random Forest', 'Extra Trees', 'SVM',
        'Gradient Boosted', 'Baseline']
X_train = X_train.drop('G3', axis='columns')
X_test = X_test.drop('G3', axis='columns')

# Instantiate the models
model1 = LinearRegression()
model2 = ElasticNet(alpha=1.0, l1_ratio=0.5)
model3 = RandomForestRegressor(n_estimators=100)
model4 = ExtraTreesRegressor(n_estimators=100)
model5 = SVR(kernel='rbf', degree=3, C=1.0, gamma='auto')
model6 = GradientBoostingRegressor(n_estimators=50)

# Dataframe for results
results = pd.DataFrame(columns=['mae', 'rmse'], index = model_name_list)

# Train and predict with each model
for i, model in enumerate([model1, model2, model3, model4, model5, model6]):
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)

    # Metrics
    mae = np.mean(abs(predictions - y_test))
    rmse = np.sqrt(np.mean((predictions - y_test) ** 2))

    # Insert results into the dataframe
    model_name = model_name_list[i]
    results.loc[model_name, :] = [mae, rmse]

# Median Value Baseline Metrics
baseline = np.median(y_train)
baseline_mae = np.mean(abs(baseline - y_test))
baseline_rmse = np.sqrt(np.mean((baseline - y_test) ** 2))

results.loc['Baseline', :] = [baseline_mae, baseline_rmse]

return results

```

```

results = evaluate(X_train, X_test, y_train, y_test)
results

```

	mae	rmse
Linear Regression	3.48512	4.4326
ElasticNet Regression	3.60805	4.57327
Random Forest	3.72601	4.61621
Extra Trees	3.7797	4.77882
SVM	3.54927	4.58147

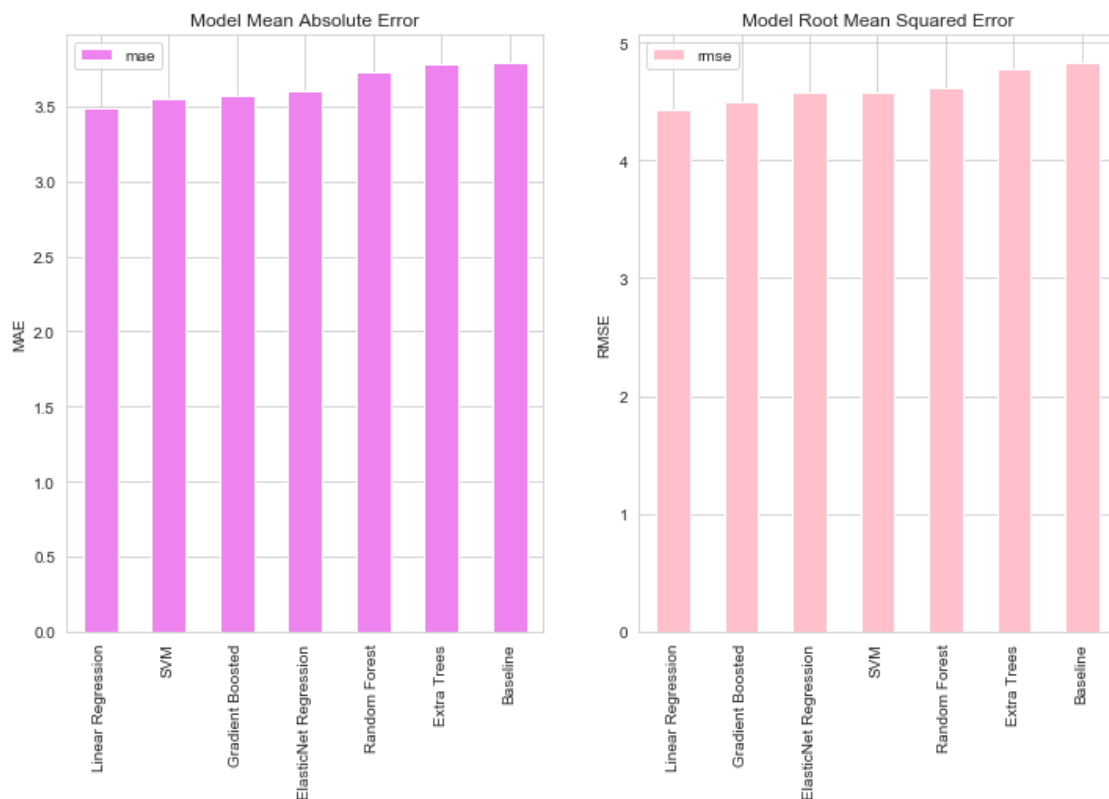
Gradient Boosted	3.57244	4.50059
Baseline	3.78788	4.82523

```
plt.figure(figsize=(12, 7))

# Root mean squared error
ax = plt.subplot(1, 2, 1)
results.sort_values('mae', ascending = True).plot.bar(y = 'mae', color = 'violet', ax = ax)
plt.title('Model Mean Absolute Error')
plt.ylabel('MAE')

# Median absolute percentage error
ax = plt.subplot(1, 2, 2)
results.sort_values('rmse', ascending = True).plot.bar(y = 'rmse', color = 'pink', ax = ax)
plt.title('Model Root Mean Squared Error')
plt.ylabel('RMSE')

plt.show()
```



Conclusion: As we see both Model Mean Absolute Error & Model Root Mean Squared Error that

the linear regression is performing the best in both cases